

## **Chapter 3: Real-time operating system (RTOS)**

After completing this chapter, you will understand what an RTOS is and how to use one. We will cover the Infineon RTOS abstraction layer and FreeRTOS but the same concepts apply to any RTOS.

### **Table of contents**

<b>3.1</b>	<b>Introduction to RTOS .....</b>	<b>3</b>
3.1.1	General .....	3
3.1.2	Tasks/threads.....	3
3.1.3	Potential pitfalls.....	3
3.1.4	Queues.....	4
3.1.5	Semaphores .....	4
3.1.6	Mutexes.....	4
3.1.7	Events/notifications.....	4
3.1.8	Interrupts.....	5
3.1.9	Timers.....	5
<b>3.2</b>	<b>RTOS abstraction layer .....</b>	<b>5</b>
<b>3.3</b>	<b>FreeRTOS .....</b>	<b>6</b>
3.3.1	Getting started .....	6
3.3.2	FreeRTOSConfig.h .....	7
3.3.3	Configuration settings .....	8
3.3.4	Kernel interface functions .....	9
3.3.5	Tasks.....	11
3.3.6	Queues.....	12
3.3.7	Semaphores and mutexes .....	13
3.3.8	Task notifications.....	16
3.3.9	Interrupts.....	17
3.3.10	Timers.....	18
3.3.11	Debugging .....	19
3.3.12	References .....	19
<b>3.4</b>	<b>Exercises .....</b>	<b>20</b>
	Exercise 1: Use a semaphore .....	20
	Exercise 2: Use a mutex .....	21
	Exercise 3: Use a queue .....	23

## Document conventions

Convention	Usage	Example
Courier New	Displays code and text commands	CY_ISR_PROTO(MyISR) ; make build
<i>Italics</i>	Displays file names and paths	<i>sourcefile.hex</i>
[bracketed, bold]	Displays keyboard commands in procedures	[Enter] or [Ctrl] [C]
Menu > Selection	Represents menu paths	File > New Project > Clone
<b>Bold</b>	Displays GUI commands, menu paths and selections, and icon names in procedures	Click the <b>Debugger</b> icon, and then click <b>Next</b> .

## 3.1 Introduction to RTOS

### 3.1.1 General

The purpose of an RTOS is to reduce the complexity of writing embedded firmware that has multiple asynchronous, response-time-critical tasks that have overlapping resource requirements. For example, you might have a device that is reading and writing data to/from a connected network, an external filesystem, and peripherals. Making sure that you deal with the timing requirement of responding to network requests while continuing to support the peripherals can be complex and, therefore, error prone. By using an RTOS, you can separate the system functions into separate tasks (called threads or tasks) and develop them in a somewhat independent fashion.

### 3.1.2 Tasks/threads

A thread/task is a specific execution context. The RTOS maintains a list of threads that are idle, halted, or running, as well as which task needs to run next (based on priority) and at what time. This function in the RTOS is called the scheduler. There are two major schemes for managing which threads/tasks/processes are active in operating systems: preemptive and co-operative.

In preemptive multitasking, the CPU completely controls which task is running and can stop and start them as required. In this scheme, the scheduler uses CPU protected modes to wrest control from active tasks, halt them, and move onto the next task. Higher priority tasks will be prioritized to run before lower priority tasks. That is, if a task is running and a higher priority task requests a turn, the RTOS will suspend the lower priority task and switch to the higher priority task. Preemptive multitasking is the scheme that is used in Windows, Linux etc.

In co-operative multitasking, each process must be a good citizen and yield control back to the RTOS. There are several mechanisms for yielding control such as RTOS delay functions, semaphores, mutexes, and queues (which we will discuss later in this document).

Most embedded RTOSes are configured in a hybrid mode. It is preemptive for higher priority tasks and "round robin" for tasks of equal priority. Higher priority tasks will always run at the expense of lower priority tasks, so it is important to yield control to give lower priority tasks a turn. If not, tasks that don't yield control will prevent lower or equal priority tasks from running at all. It is good practice to have some form of yield control mechanism in every thread to prevent such situations.

### 3.1.3 Potential pitfalls

All of this sounds great, but there are three serious bugs which can easily be created in these types of systems and these bugs can be very hard to find. These bugs are all caused by side effects of interactions between the threads. The big three are:

- Cyclic dependencies which can cause deadlocks.
- Resource conflicts when sharing memory and sharing peripherals which can cause erratic non-deterministic behavior.
- Difficulties in executing inter-process communication.

Every RTOS gives you mechanisms to deal with these problems, specifically semaphores, mutexes, and queues which we will discuss next.

### 3.1.4 Queues

A queue is a thread-safe mechanism to send data to another thread. The queue is a FIFO; you read from the front and you write to the back. If you try to read a queue that is empty your thread will suspend until something is written into it. The payload in a queue (size of each entry) and the size of the queue (number of entries) is user configurable at queue creation time. To add data to a queue, you will do a "push" or "send" operation. To pull data from a queue, you use either "pop" or "receive".

### 3.1.5 Semaphores

A semaphore is a signaling mechanism between threads. The name semaphore (originally sailing ship signal flags) was applied to computers by Dijkstra in a paper about synchronizing sequential processes. Semaphores can either be binary or counting. In the case of a counting semaphore, they are usually implemented as a simple unsigned integer. When you "set" or "give" a counting semaphore it increments the value of the semaphore. When you "get" or "take" a semaphore it decrements the value, but if the value is 0 the thread will SUSPEND itself until the semaphore is set. So, you can use a semaphore to signal between threads that something is ready. For instance, you could have a `collectDataThread` that reads data from a sensor and a `sendData` thread that sends the data up to the cloud. The `sendData` thread would "get" the semaphore which will suspend the thread UNTIL the `collectDataThread` "sets" the semaphore when it has new data available that needs to be sent.

With a binary semaphore, no matter how many times the set/give function is called, there can only be one get/take event until the semaphore is cleared. With a counting semaphore, there can be as many get/take events as there have been set/give events. These can both be useful in different situations.

### 3.1.6 Mutexes

Mutex is an abbreviation for "Mutual Exclusion." A mutex is a lock on a specific resource. If you request a mutex on a resource that is already locked by another thread, then your thread will go to sleep until the lock is released. Without a mutex, you may see strange behavior if two threads try to access the same shared resource at the same time or when one thread starts to use a shared resource but is then preempted by another thread that uses the same resource before the first thread is done. Shared resources can be things like a UART or an I2C interface, a shared block of memory, etc.

### 3.1.7 Events/notifications

Events and notifications can be likened to that of a notice board, where multiple tasks have visibility to it. Tasks can check the values in the event or notification objects to determine if a certain event/notification happened and take appropriate action. Once an event is no longer required in the system, then it can be cleared.

### 3.1.8 Interrupts

An interrupt is a signal to the processor that an event has occurred, and that immediate attention is required. An interrupt is handled with an interrupt service routine (ISR). An interrupt is a function, not a thread, so it should do what it needs to do and then exit (i.e. no infinite loop). It is also important not to call any long blocking functions inside of an ISR since it will prevent other tasks from running until it completes.

Interrupts have priorities just like threads. A higher priority interrupt can cause an ISR for a lower priority interrupt to suspend until the higher priority ISR completes. This is called a nested interrupt.

### 3.1.9 Timers

An RTOS timer allows you to schedule a function to run at a specified interval (e.g. send your data to the cloud every 10 seconds). When you set up a timer you specify the function you want run and how often you want it run.

*Note: There is a single execution of the function every time the timer expires rather than a continually executing thread, so the function should NOT have an infinite loop – it should just run and exit each time the timer calls it.*

Since the timer is a function, not a thread, make sure you don't exit the main application thread if your project has no other active threads because that will crash the system.

## 3.2 RTOS abstraction layer

In order to help make middleware as portable as possible within ModusToolbox™ applications, we have developed a thin RTOS Abstraction library. This library provides a unified API around the actual RTOS. This allows middleware libraries to be written independent of the RTOS actually selected for the application. The abstraction layer provides access to all the standard RTOS resources described above.

While the primary purpose of the library is for middleware, the abstraction layer can just as easily be used by the application code. However, since the application code generally knows what RTOS is being used, this is typically an unnecessary overhead. Also, because the abstraction only implements sufficient functionality for the middleware it is likely that many features of the RTOS are not enabled, resulting in applications that use a confusing mix of APIs.

The RTOS abstraction layer functions generally all work the same way. The basic process is:

1. Include the `cyabs_rtos.h` header file so that you have access to the RTOS functions.
2. Declare a variable of the right type (for example, `cy_mutex_t`)
3. Call the appropriate create or initialize function (for example, `cy_rtos_init_mutex`). Provide it with a reference to the variable that was created in step 2.
4. Access the RTOS object using one of the access functions (for example, `cy_rtos_set_mutex`).
5. If you don't need it anymore, free up the pointer with the appropriate de-init function (for example, `cy_rtos_deinit_mutex`).

*Note: All these functions need a pointer, so it is generally best to declare these "shared" resources as static global variables within the file that they are used. You can find the full documentation for the RTOS Abstraction layer on GitHub at:*  
<https://github.com/Infineon/abstraction-rtos>

## 3.3 FreeRTOS

Pulled from the FreeRTOS website:

"Developed in partnership with the world's leading chip companies over a 15-year period, and now downloaded every 175 seconds, FreeRTOS is a market-leading real-time operating system (RTOS) for microcontrollers and small microprocessors. Distributed freely under the MIT open source license, FreeRTOS includes a kernel and a growing set of libraries suitable for use across all industry sectors. FreeRTOS is built with an emphasis on reliability and ease of use."

Basically, FreeRTOS has been around for a long time and lots of customers have used it and are already comfortable with it. All APIs are documented online at <https://www.freertos.org/a00106.html>.

### 3.3.1 Getting started

To create a FreeRTOS application you need to do the following:

1. Create an application to start with. This is the application that you will add FreeRTOS into.
2. Add the FreeRTOS library to the application. You can do this with the Library Manager. This will also add the dependencies `clib-support` and `abstraction-rtos`.
3. Add `FREERTOS` and `RTOS_AWARE` to the `COMPONENTS` variable in the *Makefile*.

*Note:* Separate values in the list of `COMPONENTS` using spaces.

*Note:* `RTOS_AWARE` allows HAL functions that require a delay to use the RTOS delay function to allow other tasks to run instead of using a busy wait loop. Alternately, you can add `CY_RTOS_AWARE` to the `DEFINES` variable in the *Makefile*. Either one does the same thing.

4. Copy the *FreeRTOSConfig.h* file from the following directory into the root directory of your project:

```
../mtb_shared/freertos/<version>/Source/portable/COMPONENT_<core>
```

This is a template containing settings to customize how FreeRTOS behaves. More on that in a minute. The value of `<core>` should match the CPU you are using. For example, for the PSoC™ CM4 core, you would use the file from `COMPONENT_CM4`. You will typically find this specified in the BSP Makefile (`<bsp-name>.mk`) as the variable `CORE`.

5. Open the copied *FreeRTOSConfig.h* file. Remove the `#warning` line and make any other setting changes required for your application.
6. In the *main.c* file:
  - Add the required header files:

```
#include "FreeRTOS.h"
#include "task.h"
```

- Create functions for any tasks that your application requires. For example:

```
void MyTask(void *arg)
{
    (void)arg;
    /* Initialize hardware */
    for (;;)
    {
        /* Do something */
        MyActions();

        /* Allow other tasks to execute */
        vTaskDelay(500);
    }
}
```

- Setup and start any RTOS objects such as queues and semaphores.
- Create a task for each of your functions and start the RTOS scheduler:

```
int main(void)
{
    /* Initialize the device and board peripherals */
    if (CY_RSLT_SUCCESS != cybsp_init())
    {
        /* Unable to initialize BSP so HALT */
        CY_ASSERT(0);
    }

    if (pdPASS == xTaskCreate(MyTask, // Task function
                             "Task Name", // Task name
                             1024, // Task stack size
                             NULL, // Parameters passed to task
                             2, // Task priority
                             NULL) // Task handle
    )
    {
        vTaskStartScheduler();
    }

    /* vTaskStartScheduler should never return.
       If we get here, then there was a serious error */
    CY_ASSERT(0);
}
```

### 3.3.2 FreeRTOSConfig.h

This file is where you configure how FreeRTOS will work. It also contains the system specific configuration. A predefined template you can modify is in the `../mtb_shared/freertos/Source/portable/COMPONENT_<core>` directory where `<core>` is the type of CPU core. For example, if you are using the PSoC™ CM4 core, the file would be in `COMPONENT_CM4`. When you build this template, it will generate a warning. You should make a copy of this file, save it in your project's root directory and remove the warning. This file controls the APIs you can call as well as other FreeRTOS settings. For many applications, the template will work just fine as-is.

### 3.3.3 Configuration settings

The FreeRTOS [Customization](#) documentation describes the configuration options available in the *FreeRTOSConfig.h* file. The following configuration values are specific to a CM4 MCU FreeRTOS port (used for PSoC™ 6 products). They may be different for other product families.

#### 3.3.3.1 configCPU\_CLOCK\_HZ

This parameter passes the frequency of the MCU core clock required for FreeRTOS system timer configuration. The pre-configured *FreeRTOSConfig.h* file provided with the FreeRTOS library sets the `configCPU_CLOCK_HZ` value as `SystemCoreClock`, which is calculated by the MCU system startup code. You should not need to modify this parameter.

#### 3.3.3.2 configMAX\_SYSCALL\_INTERRUPT\_PRIORITY

This parameter sets the highest interrupt priority from which to call interrupt-safe FreeRTOS functions. Calling FreeRTOS functions from an interrupt with a priority higher than `configMAX_SYSCALL_INTERRUPT_PRIORITY` causes FreeRTOS to generate an exception. To avoid this exception, you can do one of the following:

- Reduce all interrupt priorities to `configMAX_SYSCALL_INTERRUPT_PRIORITY` or lower.
- Trigger an interrupt with a priority less or equal to `configMAX_SYSCALL_INTERRUPT_PRIORITY` and call the FreeRTOS functions from this interrupt handler.
- Call the FreeRTOS functions from the `traceTASKSWITCHEDOUT` macro. [See this post on FreeRTOS support forums](#).

*Note:* For the PSoC™ 6 product family, code that writes to the flash block (including the Flash/Bluetooth® Low Energy (LE)/Emulated EEPROM/DFU) uses system calls and thus the system pipe (IPC) interrupt to do so. If the IPC interrupt priority is less than or equal to `configMAX_SYSCALL_INTERRUPT_PRIORITY` and another task is running at the same or higher priority than the IPC or if a critical section is entered during the flash write this can cause issues with both blocking and partially blocking flash writes. See the [PSoC™ 6 PDL Flash Driver documentation](#) for more details.

#### 3.3.3.3 configHEAP\_ALLOCATION\_SCHEME

This parameter specifies the memory management scheme. The FreeRTOS kernel requires RAM for each created task, queue, mutex, software timer, semaphore, or event group. To manage used memory, FreeRTOS implements several different memory management schemes so that you can choose the most suitable for your application.

Each scheme is documented in the [FreeRTOS Memory Management](#) topic. The memory management implementation files are in the `freertos/Source/portable/MemMang` directory.

The available memory management schemes:

- `HEAP_ALLOCATION_TYPE1` ([heap\\_1](#)) - The simplest, does not permit memory to be freed.
- `HEAP_ALLOCATION_TYPE2` ([heap\\_2](#)) - Permits memory to be freed, but not does coalesce adjacent free blocks.
- `HEAP_ALLOCATION_TYPE3` ([heap\\_3](#)) - Simply wraps a standard malloc and free for task safety.



- `HEAP_ALLOCATION_TYPE4` ([heap\\_4](#)) - Coalesces adjacent free blocks to avoid fragmentation. Includes the absolute address-placement option.
- `HEAP_ALLOCATION_TYPE5` ([heap\\_5](#)) - As per `heap_4`, with the ability to span the heap across multiple non-adjacent memory areas.
- `NO_HEAP_ALLOCATION` - Disables the memory management, used for applications with static memory allocation.

In the `heap_3` memory scheme, your linker file must specify a sufficient size of heap and stack, and your firmware must implement and use `malloc` and `free` to allocate and release memory.

In the other memory schemes, the RTOS itself allocates a stack and heap. For these schemes, the stack defined in the Board Support Package (BSP) linker file is used only by the `main` function and the functions it calls.

#### 3.3.3.4 **configTOTAL\_HEAP\_SIZE**

This parameter specifies a total amount of RAM available for the FreeRTOS heap. This parameter ignored when `heap_3` memory scheme is used.

#### 3.3.3.5 **configMINIMAL\_STACK\_SIZE**

This parameter specifies the size of the stack used by the idle task. It is not recommended to reduce the default parameter value.

### 3.3.4 **Kernel interface functions**

While the [API lists a dozen kernel functions](#), there are only 6 needed in day to day use, and 4 of them are paired sets.

The first function of interest is `vTaskStartScheduler` which is what gets the RTOS scheduler going. It is usually called at the end of `main` and it should never return.

The others are broken down by what type of task scheduler you're using. The `FreeRTOSConfig.h` file controls the scheduler type; see the FreeRTOS [Customization](#) documentation.

#### 3.3.4.1 **Cooperative**

- `taskYIELD` - Lets the kernel schedule someone else. Note this is not like a delay which blocks the task for a specified period of time. Instead it allows tasks at the same or higher priority to run. Lower priority tasks will never run because of a yield.

*Note: This can be used in a preemptive system if there are specific places where you want to let other tasks run, instead of waiting for the next system tick to kick you out. For example, you may want to yield before starting a critical section.*

If you want to allow any task to run, even a lower priority task, you can use `vTaskDelay`. We'll cover that in just a minute.

### 3.3.4.2 Preemptive

There are critical section functions that are useful in a preemptive system to prevent a task from being preempted for some critical operation.

*Note: FreeRTOS API functions must not be called in a critical section.*

- `taskENTER_CRITICAL/taskEXIT_CRITICAL` – Defines a block of code that must be performed atomically. Must not be called from an ISR.
- `taskENTER_CRITICAL_FROM_ISR/taskEXIT_CRITICAL_FROM_ISR` – Defines a block of code that must be performed atomically. Must only be called from an ISR.

In general, you should stay in the same scope for the enter critical section/exit critical section. In some cases, you may want to call other APIs and if those APIs do not rely on interrupts you're fine. That said, like interrupts, critical sections should be kept short whenever possible.

### 3.3.5 Tasks

Documentation on task creation can be found at: <https://www.freertos.org/a00019.html>

To create a task, you use the [xTaskCreate](#) API. Its prototype is:

```
BaseType_t xTaskCreate( TaskFunction_t pxTaskCode,  
                        const char * const pcName,  
                        const configSTACK_DEPTH_TYPE usStackDepth,  
                        void * const pvParameters,  
                        UBaseType_t uxPriority,  
                        TaskHandle_t * const pxCreatedTask )
```

Of particular interest to us is the priority. Zero is the lowest and generally reserved for the idle task which may or may not feed a watchdog timer (WDT). In a preemptive system there are (generally) two basic rules at play:

- Task with highest priority will run first, and never give up the CPU until the task is blocked
- If 2 or more tasks with the same priority do not give up the CPU (they don't block), then FreeRTOS will share the CPU between them. *FreeRTOSConfig.h* determines how this sharing happens.

You can make a task give up the CPU (that is, block it) by calling:

<a href="#">vTaskDelay</a>	This simply blocks the task so that something else can run; you decide how long you want it to block (in ticks)
----------------------------	---

*Note:* You can use the function *pdMS\_TO\_TICKS* to milliseconds to ticks. For example, *pdMS\_TO\_TICKS(500)* will give the correct number of ticks for a 500 ms delay. Usually the tick is set to 1 ms.

<a href="#">xQueueSend</a>	If the Queue you are sending to is full, this task will be blocked until there is space in the queue or until the timeout elapses.
<a href="#">xQueueReceive</a>	If the Queue you are reading from is empty, this task will be blocked until there is data in the queue or until the timeout elapses.
<a href="#">xSemaphoreTake</a>	The task will be blocked if the semaphore has not been given or until the timeout elapses.

There are other ways for tasks to become blocked as well, but these are the most common.

### 3.3.6 Queues

You will use this in [Exercise 3](#):

Documentation on queues can be found at <https://www.freertos.org/a00018.html>

You can communicate between tasks by using Queues or Semaphores. In order to use the queue functions, you must include *queue.h* in your source code. Let's create an example to communicate between two tasks. Here's some sample code:

```
#include "FreeRTOS.h"
#include "task.h"
#include "queue.h"
#include "stdio.h"

// Global Queue Handle
static QueueHandle_t qh = 0;

void task_tx(void *p)
{
    int myCount = 0;
    while (1)
    {
        myCount++;
        if (!xQueueSend(qh, &myCount, pdMS_TO_TICKS(500)))
        {
            puts("Failed to send item to queue within 500ms");
        }
        vTaskDelay(1000);
    }
}

void task_rx(void *p)
{
    int myValue = 0;
    while (1)
    {
        if (!xQueueReceive(qh, &myValue, pdMS_TO_TICKS(10000)))
        {
            puts("Failed to receive item within 10000 ms");
        }
        else
        {
            printf("Received: %u\n", myValue);
        }
    }
}

int main(void)
{
    qh = xQueueCreate(1, sizeof(int));

    xTaskCreate(task_tx, (char *)"t1", 2048, 0, 1, 0);
    xTaskCreate(task_rx, (char *)"t2", 2048, 0, 1, 0);
    vTaskStartScheduler();

    CY_ASSERT(0); // Should never get here
}
```

*Note:*

- In `main`, we create the Queue before creating tasks, because sending to an un-initialized Queue will crash the system.
- In `task_tx`, we send one item every second. If the queue is full the task will block for 500ms (set by the timeout value). The value is specified in ticks so the macro `pdMS_TO_TICKS` is required to convert to milliseconds. Once the queue is no longer full, the item will be sent. If the queue stays full for 500ms, the thread unblocks due to the timeout and we print a failure message.
- In `task_rx`, we receive one item, and we do not use `vTaskDelay`. This is because if there is nothing in the queue, FreeRTOS will block the task from running. `xQueueReceive` will unblock the task as soon as an item shows up in the queue. If nothing shows up in 10000ms, the task will unblock anyway - that is set by the value of 10000 for the timeout in the `xQueueReceive` function call.
- If the priority of the receiving queue(`task_rx`) is higher, FreeRTOS will switch tasks the moment `xQueueSend` happens, and the next line inside `task_tx` will not execute since CPU will be switched over to `task_rx`.

### 3.3.7 Semaphores and mutexes

You will use this in [Exercise 1](#): and [Exercise 2](#):

Documentation on semaphores and mutexes can be found at: <https://www.freertos.org/a00113.html>

In FreeRTOS, Mutex functions are grouped with Semaphores since a mutex is just a special case of a semaphore. In fact, there are functions to create a mutex, but you just use semaphore functions to lock and unlock a mutex. Let me show you with an example. Note that you need to include *semphr.h* in your source code to use semaphores (or mutexes).

### 3.3.7.1 Mutexes

Semaphores in FreeRTOS can be used as mutexes with a [priority inversion](#) mechanism such that a lower priority task that holds a mutex for a shared resource will prevent a higher priority task from running if that higher priority task requires the same mutex. One of the best examples of a mutex is to guard a resource or a door with a key. For instance, let's say you have an SPI BUS, and only one task should use it at a time.

A mutex will only allow ONE task to get past an `xSemaphoreTake` operation and other tasks will be put to sleep if they reach this function at the same time. In this example, the two tasks use a semaphore as a mutex.

```
#include "FreeRTOS.h"
#include "task.h"
#include "semphr.h"

static SemaphoreHandle_t spi_bus_lock=0;

void task_one()
{
    while (1)
    {
        if (xSemaphoreTake(spi_bus_lock, 1000))
        {
            // Use Guarded Resource

            // Give Semaphore back:
            xSemaphoreGive(spi_bus_lock);
        }
    }
}

void task_two()
{
    while (1)
    {
        if (xSemaphoreTake(spi_bus_lock, 1000))
        {
            // Use Guarded Resource

            // Give Semaphore back:
            xSemaphoreGive(spi_bus_lock);
        }
    }
}

int main(void)
{
    spi_bus_lock = xSemaphoreCreateMutex();
    xTaskCreate(task_one, (char *) "t1", 2048, 0, 1, 0);
    xTaskCreate(task_two, (char *) "t2", 2048, 0, 1, 0);
    vTaskStartScheduler();
    CY_ASSERT(0);
}
```

Note that the semaphore was created using the function `xSemaphoreCreateMutex` which takes care of setting up the semaphore with priority inversion.

### 3.3.7.2 Binary semaphore

A binary semaphore can also be used like a mutex, but binary semaphores doesn't provide a priority inversion mechanism. Binary semaphores are better suited for helper tasks for interrupts. For example, if you have an interrupt and you don't want to do a lot of processing inside the interrupt, you can use a helper task. To accomplish this, you can perform a semaphore "give" operation inside the interrupt, and a dedicated task will block on the corresponding take operation.

```
#include "FreeRTOS.h"
#include "task.h"
#include "semphr.h"

#define LONG_TIME 0xffff

static SemaphoreHandle_t event_signal = 0;

void System_Interrupt()
{
    static BaseType_t xHigherPriorityTaskWoken;
    xSemaphoreGiveFromISR(event_signal, &xHigherPriorityTaskWoken);

    /* Yield current task if a higher priority task is now unblocked */
    portYIELD_FROM_ISR(xHigherPriorityTaskWoken);
}

void system_interrupt_task()
{
    while (1)
    {
        if (pdTRUE == xSemaphoreTake(event_signal, LONG_TIME))
        {
            // Process the interrupt
        }
    }
}

int main(void)
{
    vSemaphoreCreateBinary(event_signal); // Create the semaphore
    xSemaphoreTake(event_signal, 0);      // Take semaphore after creating it.
    /* create your task, assuming it has an interrupt to give the semaphore*/
    vTaskStartScheduler();
    CY_ASSERT(0);
}
```

The above code shows an example of deferred interrupt processing. The idea is that you don't want to process the interrupt inside `System_Interrupt` because you'd be in a critical section with system interrupts globally disabled, therefore, you can potentially lock up the system or destroy real-time processing if the interrupt processing takes too long.

In the example, the ISR function `System_Interrupt` just gives the semaphore (with the `FromISR` designation) and then exits. The actual processing happens inside the function `system_interrupt_task`, which waits for the semaphore and then does the processing. This way, other critical tasks that have a higher priority than `system_interrupt_task` can run first.

Another way to use binary semaphore is to wake up one task from another task by giving the semaphore. The semaphore will essentially act like a signal that may indicate: "Something happened, now go do the work in another task".

Again, remember that when you create a semaphore in FreeRTOS, it is ready to be taken, so you may want to take the semaphore after you create it such that the task waiting on this semaphore will block until given by somebody.

### 3.3.7.3 Counting semaphore

As you learned earlier, a counting semaphore keeps track of how many times it was given rather than being a binary event. With a counting semaphore, you specify the depth and the initial count when you create the semaphore. Its function prototype looks like this:

```
SemaphoreHandle_t xSemaphoreCreateCounting( UBaseType_t uxMaxCount,  
                                             UBaseType_t uxInitialCount )
```

### 3.3.8 Task notifications

Every FreeRTOS task has a 32-bit notification value. An RTOS task notification is an event sent directly to a task that can unblock the receiving task, and optionally update the receiving task's notification value. Unblocking an RTOS task with a direct notification is 45% faster and uses less RAM than unblocking a task with a binary semaphore.

Notifications are sent using the `xTaskNotify` and `xTaskNotifyGive` functions and received using `xTaskNotifyWait` and `ulTaskNotifyTake` functions. A task that calls one of the receiving functions will block for a specified amount of time until another task calls one of the sending functions, specifying the blocked receiving task.

Task notifications can update the receiving task's notification value in several ways:

- Set the receiving task's notification value with or without overwriting previous notification values
- Set one or more bits in the receiving task's notification value
- Increment the receiving task's notification value

If you want to use a notification as a binary semaphore you should use the `xTaskNotifyGive` and the `ulTaskNotifyTake` functions. If you want to use a notification to deliver more data than a simple binary semaphore provides you should use the `xTaskNotify` and the `xTaskNotifyWait` functions.

**Note:** *The functions mentioned thus far for sending task notifications are not ISR safe! If you want to send a notification from an ISR you will need to call the ISR safe equivalent of the function you want:*

*`xTaskNotifyFromISR` or `vTaskNotifyGiveFromISR`*

There is a complete description of, as well as a plethora of usage examples for notifications, available from the [FreeRTOS's Task Notification](#) section of the website.



### 3.3.9 Interrupts

Interrupts will do just as you would expect; stop the current task so that they can be serviced. In fact, this is how the preemptive/time slice systems are implemented, with a system tick interrupt at the lowest interrupt priority. Here is an ISR from a CAPSENSE™ code example:

```
static void capsense_timer_callback(TimerHandle_t xTimer)
{
    capsense_command_t cmd = CAPSENSE_SCAN;
    BaseType_t xYieldRequired;

    (void)xTimer;

    /* Send command to start CapSense scan */
    xQueueSendFromISR(capsense_cmd_q, &cmd, &xYieldRequired);

    /* Yield current task if a higher priority task is now unblocked */
    portYIELD_FROM_ISR(xYieldRequired);
}
```

You may be wondering what `portYIELD_FROM_ISR` does. Well let's imagine task 1 with a priority of 2 is executing when the ISR comes in. Provided interrupts are enabled and task 1 is not in a critical section, the system will switch to the ISR and run it.

Notice that the ISR calls the function `xQueueSendFromISR` that may unblock a task. Also notice the function has the "FromISR" suffix so that it is safe to call inside the ISR. If this unblocks a task with a higher priority than the one which was executing when the ISR fired, `xYieldRequired` will be set to true. Continuing the example, assume `xQueueSendFromISR` unblocks task 2 and it has a priority of 3, so `xYieldRequired` is true.

When `portYIELD_FROM_ISR` is called with a value of true, this queues up a scheduler change as the next action once the ISR is complete. From the user's point of view, the system processing is task 1 → ISR → task 2. Once task 2 is done (or blocked again), the scheduler will go back to task 1.

If `portYIELD_FROM_ISR` is false, the task will be executed once the ISR exits since nothing else of a higher priority asked for a turn. From the user's perspective it looks like task 1 → ISR → task 1. Failing to call `portYIELD_FROM_ISR` has the same behavior as calling it with false.

With our example if you failed to call `portYIELD_FROM_ISR`, the behavior the user will observe is task 1 → ISR → task 1 → task 2. A scheduler event like an ISR or system tick is required before task 1 is suspended, so this would delay when the higher priority task (task 2 in this case) gets to run.

### 3.3.10 Timers

Timers can be set up to only fire once or to fire periodically. Each time a timer expires, it will call a function that you registered when you created the timer. To use a timer:

- `#include "timers.h"` in your C file.
- Create a global variable of type `TimerHandle_t`.
- Call `xTimerCreate`.
- Call `xTimerStart`.
- Create the callback function.

For example, the following will setup and start a timer that expires every 1000 ms:

```
#include "FreeRTOS.h"
#include "task.h"
#include "timers.h"

static void timerCallback(TimerHandle_t myTimer);

TimerHandle_t myTimer;

void timerCallback(TimerHandle_t myTimer)
{
    /* Functionality each time the timer expires goes here */
}

int main(void)
{
    myTimer = xTimerCreate("myTimer", pdMS_TO_TICKS(1000), pdTRUE, (void *) 0,
                           timerCallback);

    xTimerStart(myTimer, 0);

    /* Create any tasks that your application requires here */

    vTaskStartScheduler();
    CY_ASSERT(0);
}
```

Documentation on timers can be found at <https://www.freertos.org/FreeRTOS-Software-Timer-API-Functions.html>.

### 3.3.11 Debugging

In order for RTOS-aware debugging to work properly using OpenOCD, the following items should be included in the code:

Global variable:

```
volatile int uxTopUsedPriority;
```

At the top of main:

```
uxTopUsedPriority = configMAX_PRIORITIES - 1;
```

### 3.3.12 References

The information above comes for a variety of sources: Some of the following sources have more recommendations/information.

- <https://www.freertos.org/a00106.html>
- <https://github.com/Infineon/freertos>
- [http://socialledge.com/sjsu/index.php/FreeRTOS\\_Tutorial](http://socialledge.com/sjsu/index.php/FreeRTOS_Tutorial)

## 3.4 Exercises

In the following exercises, the CY8CKIT-062S2-43012 kit is used but feel free to use any available BSP to explore the family or solution of your choice. The exact steps and starting application that you use may be different than what's shown here, but the concepts will be the same.

*Note: GPIO functions (to control LEDs and monitor buttons) and interrupts are used in these exercises. These functions are covered in more detail in the level 2 MCU classes. You can refer to that material if you need additional explanation on how those functions work.*

### Exercise 1: Use a semaphore

In this exercise you will add FreeRTOS to an empty application and use a binary semaphore to signal between a button interrupt and thread that controls an LED. This material is covered in [3.3.7](#).

- ☐ 1. Use Project Creator to create a new application using the "Empty App" as the template application. Name it **ch03\_ex01\_semaphore**.

- ☐ 2. Use the Library Manager to add the FreeRTOS library to the application.

*Note: This will also add the abstraction-freertos and clib-support libraries.*

- ☐ 3. Edit the *Makefile* to add `FREERTOS` and `RTOS_AWARE` to the `COMPONENTS` variable.

*Note: Separate the names of values in the COMPONENTS variable using spaces, NOT commas. For example:*  
`COMPONENTS = FREERTOS RTOS_AWARE`

- ☐ 4. Copy *FreeRTOSConfig.h* from the FreeRTOS library to the root directory of the application and edit it to remove the compiler warning.

*Note: Use the file from `../mtb_shared/freertos/<version>/Source/portable/COMPONENT_CM4`*

- ☐ 5. Add the required headers to *main.c*. Don't forget to include the semaphore header file.

- ☐ 6. Create a static global variable of type `SemaphoreHandle_t` for the button semaphore.

- ☐ 7. In main:
  - a. Initialize the User LED (`CYBSP_USER_LED`) as a strong drive output (see below).
  - b. Initialize the User Button (`CYBSP_USER_BTN`) as an input with a resistive pullup and an initial state of one. Register an interrupt callback and enable falling edge interrupt events (see below).
  - c. Create the semaphore. Take the semaphore right after you create it so that the task will start out blocked.
  - d. Create the LED task.
  - e. Start the RTOS scheduler.

The code to initialize the LED and the User Button will look like this:

```
/* Initialize the User LED */
result = cyhal_gpio_init(CYBSP_USER_LED, CYHAL_GPIO_DIR_OUTPUT,
                        CYHAL_GPIO_DRIVE_STRONG, CYBSP_LED_STATE_OFF);

/* Initialize button with an interrupt of priority 3 */
cyhal_gpio_init(CYBSP_USER_BTN, CYHAL_GPIO_DIR_INPUT, CYHAL_GPIO_DRIVE_PULLUP, 1);
cyhal_gpio_register_callback(CYBSP_USER_BTN, &button_interrupt_data);
cyhal_gpio_enable_event(CYBSP_USER_BTN, CYHAL_GPIO_IRQ_FALL, 3, true);
```



8. Write the button ISR so that it gives the semaphore.

The function for the ISR will look like this:

```
void button_interrupt_handler(void* handler_arg, cyhal_gpio_event_t event)
{
    /* ISR code */
}
```

Don't forget to use the `FromISR` version of the function and to use `portYIELD_FROM_ISR` to yield to a higher task.



9. Create a structure to configure the interrupt handler. It will look like this:

```
/* Structure for GPIO interrupt */
cyhal_gpio_callback_data_t button_interrupt_data =
{
    .callback      = button_interrupt_handler,
    .callback_arg = NULL
};
```

**Note:** *This should be declared as a global variable. It must be declared after the ISR function is declared since it must exist in order to specify it as the callback.*



10. Write the LED task so that it takes the semaphore and then toggles the LED.

The following code will toggle the LED:

```
cyhal_gpio_toggle(CYBSP_USER_LED);
```

Use `portMAX_DELAY` for the semaphore timeout so that it will wait forever.



11. Build and program. Press the button to toggle the LED.

## Exercise 2: Use a mutex

An LED may behave strangely if two threads try to blink it at the same time. In this exercise you will use a mutex to lock access to one button at a time. This material is covered in [3.3.7](#).



1. Use Project Creator to create a new application using Empty App as the template application. Name it **ch03\_ex02\_mutex**.

- ☐ 2. Add the FreeRTOS library, add `FREERTOS` and `RTOS_AWARE` to the `COMPONENTS` variable in the *Makefile*, copy *FreeRTOSConfig.h* from the library to the application's root directory and remove the warning.
- ☐ 3. Replace *main.c* with the following code (copy and paste). It creates two threads – one that blinks the user LED at 2 Hz and the other blinks the same LED at 5 Hz when the user button is pressed.

```
#include "cy_pdl.h"
#include "cyhal.h"
#include "cybsp.h"
#include "FreeRTOS.h"
#include "task.h"
#include "semphr.h"

volatile int uxTopUsedPriority;

//TODO create variable for mutex

/* LED task 1 - blink at 2 Hz */
void led_task1()
{
    while (1)
    {
        //TODO lock the mutex

        cyhal_gpio_toggle(CYBSP_USER_LED);

        //TODO unlock the mutex

        vTaskDelay (1000/2); /* 2 Hz delay and give the other task a turn */
    }
}

/* LED task 2 - blink at 5 Hz when button is pressed */
void led_task2()
{
    while (1)
    {
        //TODO lock the mutex

        while (CYBSP_BTN_PRESSED == cyhal_gpio_read(CYBSP_USER_BTN)) /* Button is pressed */
        {
            cyhal_gpio_toggle(CYBSP_USER_LED);
            vTaskDelay (1000/5);
        }

        //TODO unlock the mutex

        vTaskDelay (1); /* Give the other task a turn when button is not pressed */
    }
}

int main(void)
{
    uxTopUsedPriority = configMAX_PRIORITIES - 1;

    /* Initialize the device and board peripherals */
    cybsp_init();
    __enable_irq();

    /* Initialize the User LED */
    cyhal_gpio_init(CYBSP_USER_LED, CYHAL_GPIO_DIR_OUTPUT,
                   CYHAL_GPIO_DRIVE_STRONG, CYBSP_LED_STATE_OFF);

    /* Initialize User Button */
    cyhal_gpio_init(CYBSP_USER_BTN, CYHAL_GPIO_DIR_INPUT, CYHAL_GPIO_DRIVE_PULLUP, 1);

    //TODO Create Mutex

    /* Create tasks with a stack of 1024, no parameters, and a priority of 1 */
    xTaskCreate(led_task1, (char *)"led_task1", 1024, 0, 1, NULL);
    xTaskCreate(led_task2, (char *)"led_task2", 1024, 0, 1, NULL);

    /* Start the RTOS scheduler */
    vTaskStartScheduler();
    CY_ASSERT(0);
}
```

- ☐ 4. Program the application to your kit. Notice what happens when you hold down the user button. Does the LED blink at the desired 5 Hz?
- ☐ 5. Look at the `TODO` comments to add a mutex to the two tasks so that each one prevents the other from blinking the LED when it needs access.
- ☐ 6. Program the modified application to your kit. Hold down the button and observe the behavior with the mutex included.

### Exercise 3: Use a queue

In this exercise, you will use a queue to send a message to indicate the number of times to blink an LED. This material is covered in [3.3.6](#).

- ☐ 1. Use Project Creator to create a new application. This time, instead of using one of the standard template applications, use the **Browse** button and use the solution from exercise 1 as the template application. Name the new application **ch03\_ex03\_queue**.
- ☐ 2. Change the Semaphore to a Queue with a width of `uint_32` and a depth of 10.

*Note: Don't forget to include the header file for the Queue API.*

- ☐ 3. Add a static variable called `count` inside the button ISR that is incremented each time the button is pressed. Send the value of `count` to the queue from the ISR.

*Note: Don't forget to use the `FromISR` version of the queue send function.*

- ☐ 4. In the LED task, wait until there is data in the queue. When a new value is available, blink the LED the number of times specified by the value from the queue.
- ☐ 5. Build and program. Press the button a few times to see the LED blink count increase with each button press.

#### Trademarks

All referenced product or service names and trademarks are the property of their respective owners.

**Published by**  
**Infineon Technologies AG**  
**81726 Munich, Germany**

**© 2024 Infineon Technologies AG.**  
**All Rights Reserved.**

#### IMPORTANT NOTICE

The information given in this document shall in no event be regarded as a guarantee of conditions or characteristics ("Beschaffheitsgarantie").

With respect to any examples, hints or any typical values stated herein and/or any information regarding the application of the product, Infineon Technologies hereby disclaims any and all warranties and liabilities of any kind, including without limitation warranties of non-infringement of intellectual property rights of any third party.

In addition, any information given in this document is subject to customer's compliance with its obligations stated in this document and any applicable legal requirements, norms and standards concerning customer's products and any use of the product of Infineon Technologies in customer's applications.

The data contained in this document is exclusively intended for technically trained staff. It is the responsibility of customer's technical departments to evaluate the suitability of the product for the intended application and the completeness of the product information given in this document with respect to such application.

For further information on the product, technology, delivery terms and conditions and prices please contact your nearest Infineon Technologies office ([www.infineon.com](http://www.infineon.com)).

#### WARNINGS

Due to technical requirements products may contain dangerous substances. For information on the types in question please contact your nearest Infineon Technologies office.

Except as otherwise explicitly approved by Infineon Technologies in a written document signed by authorized representatives of Infineon Technologies, Infineon Technologies' products may not be used in any applications where a failure of the product or any consequences of the use thereof can reasonably be expected to result in personal injury.